

Lecture 3

C7B Optimization

Hilary 2011

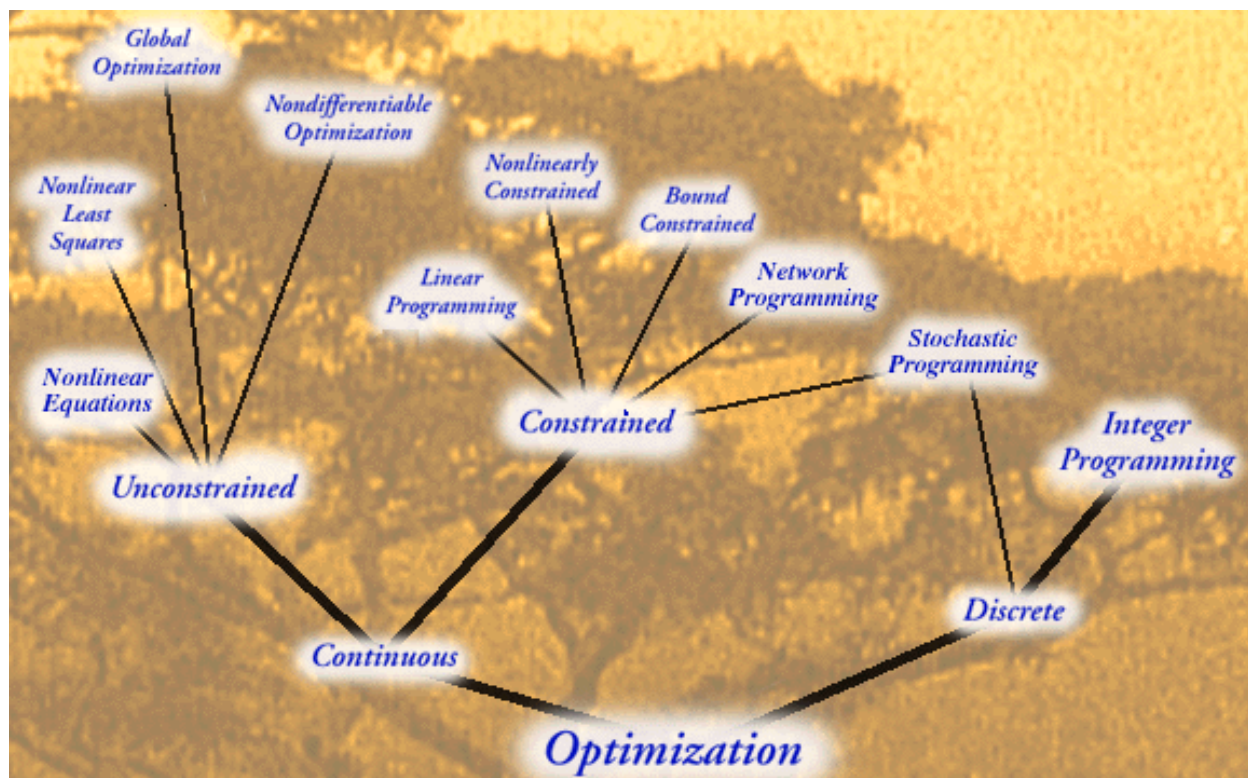
A. Zisserman

Cost functions with special structure:

- Levenberg-Marquardt algorithm
- Dynamic Programming
 - chains
 - applications

First: review Gauss-Newton approximation

The Optimization Tree



Summary of minimizations methods

Update $\mathbf{x}_{n+1} = \mathbf{x}_n + \delta\mathbf{x}$

1. Newton.

$$\mathbf{H} \delta\mathbf{x} = -\mathbf{g}$$

2. Gauss-Newton.

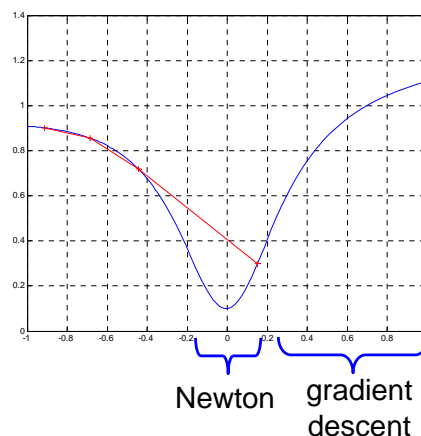
$$2\mathbf{J}^\top \mathbf{J} \delta\mathbf{x} = -\mathbf{g}$$

3. Gradient descent.

$$\lambda \delta\mathbf{x} = -\mathbf{g}$$

Levenberg-Marquardt algorithm

- Away from the minimum, in regions of negative curvature, the Gauss-Newton approximation is not very good.
- In such regions, a simple steepest-descent step is probably the best plan.
- The Levenberg-Marquardt method is a mechanism for varying between steepest-descent and Gauss-Newton steps depending on how good the $\mathbf{J}^\top \mathbf{J}$ approximation is locally.



- The method uses the modified Hessian

$$H(\mathbf{x}, \lambda) = 2J^T J + \lambda I$$

- When λ is small, H approximates the Gauss-Newton Hessian.
- When λ is large, H is close to the identity, causing steepest-descent steps to be taken.

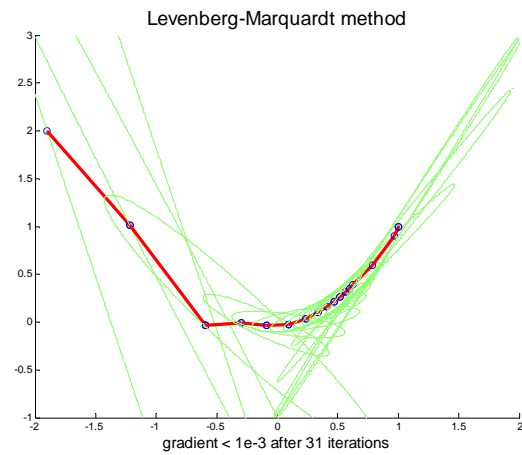

LM Algorithm

$$H(\mathbf{x}, \lambda) = 2J^T J + \lambda I$$

1. Set $\lambda = 0.001$ (say)
2. Solve $\delta \mathbf{x} = -H(\mathbf{x}, \lambda)^{-1} \mathbf{g}$
3. If $f(\mathbf{x}_n + \delta \mathbf{x}) > f(\mathbf{x}_n)$, increase λ ($\times 10$ say) and go to 2.
4. Otherwise, decrease λ ($\times 0.1$ say), let $\mathbf{x}_{n+1} = \mathbf{x}_n + \delta \mathbf{x}$, and go to 2.

Note : This algorithm does not require explicit line searches.

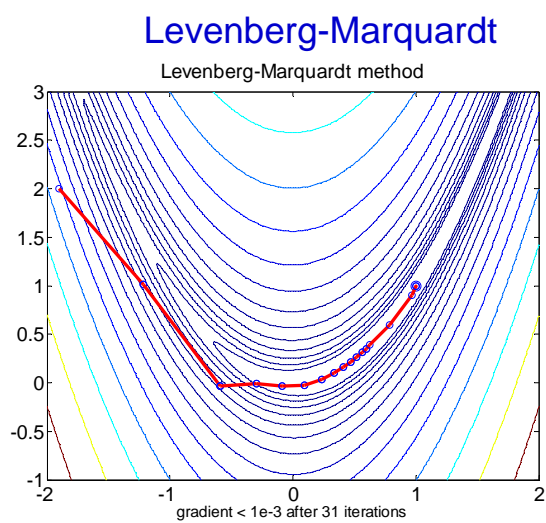
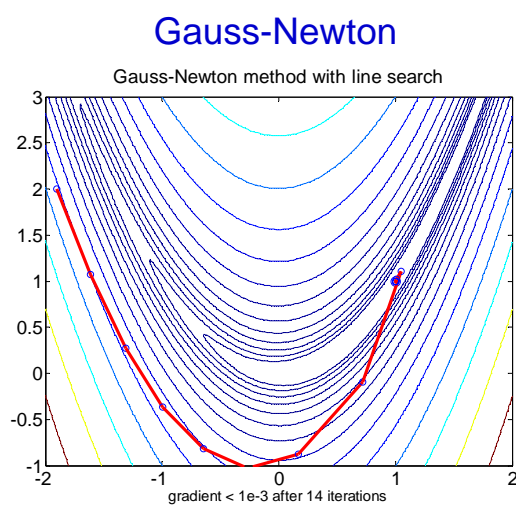
Example



- Minimization using Levenberg-Marquardt (no line search) takes 31 iterations.

Matlab: lsqnonlin

Comparison



- more iterations than Gauss-Newton, but
- no line search required,
- and more frequently converges

Case study – Bundle Adjustment (non-examinable)



Problem statement

- **Given:** n matching image points \mathbf{x}_j^i over m views
- **Find:** the cameras P^i and the 3D points \mathbf{X}_j such that $\mathbf{x}_j^i = P^i \mathbf{X}_j$

$$\min_{P^i, \mathbf{X}_j} \sum_{j \in \text{points}} \sum_{i \in \text{views}} d(\mathbf{x}_j^i, P^i \mathbf{X}_j)^2$$

Notation:

- A 3D point \mathbf{X}_j is imaged in the “ i ” th view as
- | | | |
|--------------|---|---------------------|
| P | : | 3×4 matrix |
| \mathbf{X} | : | 4-vector |
| \mathbf{x} | : | 3-vector |

$$\mathbf{x}_j^i = P^i \mathbf{X}_j$$

Number of parameters

$$\min_{P^i, \mathbf{X}_j} \sum_{j \in \text{points}} \sum_{i \in \text{views}} d(\mathbf{x}_j^i, P^i \mathbf{X}_j)^2$$

- for each camera there are 6 parameters
- for each 3D point there are 3 parameters

a total of $6m + 3n$ parameters must be estimated


- e.g. 50 frames, 1000 points: 3300 unknowns

Example

image sequence



cameras and points



Sparse form of the Jacobian matrix

- Image point \mathbf{x}_j^i does not depend on the parameters of any camera other than \mathbf{p}^i .
- Thus,

$$\partial \mathbf{x}_j^i / \partial \mathbf{p}^k = 0$$


unless $i = k$.

- Similarly, image point \mathbf{x}_j^i does not depend on any 3D point except \mathbf{X}_j .

$$\partial \mathbf{x}_j^i / \partial \mathbf{X}_k = 0$$

unless $j = k$.

Form of the Jacobian and Gauss-Newton Hessian for the bundle-adjustment problem consisting of 3 cameras and 4 points.



By taking advantage of this sparse form, one iterative update of the LM algorithm

- $H(x, \lambda) = 2J^T J + \lambda I$
- Solve $\delta x = -H(x, \lambda)^{-1} g$

can be solved in $O(N)$ rather than $O(N^3)$, where N is the total number of parameters

Application: Augmented reality

original sequence



Augmentation



Dynamic programming


- Discrete optimization
- Each variable x has a finite number of possible states
- Applies to problems that can be decomposed into a sequence of stages
- Each stage expressed in terms of results of fixed number of previous stages
- The cost function need not be convex
- The name “dynamic” is historical
- Also called the “Viterbi” algorithm

Consider a cost function $f(\mathbf{x}) : \mathbb{R}^n \rightarrow \mathbb{R}$ of the form

$$f(\mathbf{x}) = \sum_{i=1}^n m_i(x_i) + \sum_{i=2}^n \phi_i(x_{i-1}, x_i)$$

where x_i can take one of h values

e.g. $h=5, n=6$




Complexity of minimization:

- exhaustive search $O(h^n)$
- dynamic programming $O(nh^2)$

Example 1

$$f(\mathbf{x}) = \sum_{i=1}^n m_i(x_i) + \sum_{i=2}^n \phi(x_{i-1}, x_i)$$

$$f(\mathbf{x}) = \sum_{i=1}^n \underbrace{(x_i - d_i)^2}_{\text{closeness to measurements}} + \sum_{i=2}^n \underbrace{\lambda^2 (x_i - x_{i-1})^2}_{\text{smoothness}}$$



Motivation: complexity of stereo correspondence

Objective: compute horizontal displacement for matches between left and right images




x_i is spatial shift of i 'th pixel $\rightarrow h = 40$

\mathbf{x} is all pixels in row $\rightarrow n = 256$

Complexity $O(40^{256})$ vs $O(256 \times 40^2)$

$$f(\mathbf{x}) = \sum_{i=1}^n m_i(x_i) + \sum_{i=2}^n \phi(x_{i-1}, x_i)$$



Key idea: the optimization can be broken down into n sub-optimizations


Step 1: For each value of x_2 determine the best value of x_1

- Compute

$$\begin{aligned} S_2(x_2) &= \min_{x_1} \{m_2(x_2) + m_1(x_1) + \phi(x_1, x_2)\} \\ &= m_2(x_2) + \min_{x_1} \{m_1(x_1) + \phi(x_1, x_2)\} \end{aligned}$$

- Record the value of x_1 for which $S_2(x_2)$ is a minimum

To compute this minimum for all x_2 involves $O(h^2)$ operations



Step 2: For each value of x_3 determine the best value of x_2 **and** x_1

- Compute

$$S_3(x_3) = m_3(x_3) + \min_{x_2} \{S_2(x_2) + \phi(x_2, x_3)\}$$

- Record the value of x_2 for which $S_3(x_3)$ is a minimum

Again, to compute this minimum for all x_3 involves $O(h^2)$ operations

Note $S_k(x_k)$ encodes the lowest cost partial sum for all nodes up to k which have the value x_k at node k , i.e.

$$S_k(x_k) = \min_{x_1, x_2, \dots, x_{k-1}} \sum_{i=1}^k m_i(x_i) + \sum_{i=2}^k \phi(x_{i-1}, x_i)$$

Viterbi Algorithm

- Initialize $S_1(x_1) = m_1(x_1)$

- For $k = 2 : n$

$$S_k(x_k) = m_k(x_k) + \min_{x_{k-1}} \{S_{k-1}(x_{k-1}) + \phi(x_{k-1}, x_k)\}$$

$$b_k(x_k) = \arg \min_{x_{k-1}} \{S_{k-1}(x_{k-1}) + \phi(x_{k-1}, x_k)\}$$

- Terminate

$$x_n^* = \arg \min_{x_n} S_n(x_n)$$

- Backtrack

$$x_{i-1} = b_i(x_i)$$




Complexity $O(nh^2)$

Example 2

$$f(\mathbf{x}) = \sum_{i=1}^n (x_i - d_i)^2 + \sum_{i=2}^n g_{\alpha, \lambda}(x_i - x_{i-1})$$

where

$$g_{\alpha, \lambda}(\Delta) = \min(\lambda^2 \Delta^2, \alpha) = \begin{cases} \lambda^2 \Delta^2 & \text{if } |\Delta| < \sqrt{\alpha}/\lambda \\ \alpha & \text{otherwise.} \end{cases}$$



Note, $f(x)$ is not convex

Note

This type of cost function often arises in MAP estimation

$$\mathbf{x}^* = \arg \max_{\mathbf{x}} p(\mathbf{x}|\mathbf{y})$$

measurements

$$= \arg \max_{\mathbf{x}} p(\mathbf{y}|\mathbf{x})p(\mathbf{x}) \quad \text{Bayes' rule}$$

$$\sim \prod_i^n e^{-\frac{(x_i - y_i)^2}{2\sigma^2}} e^{-\beta^2(x_i - x_{i-1})^2}$$

e.g. for Gaussian measurement errors, and first order smoothness

Use negative log to obtain a cost function of the form

$$f(\mathbf{x}) = \sum_{i=1}^n \underbrace{(x_i - y_i)^2}_{\text{from likelihood}} + \sum_{i=2}^n \underbrace{\lambda^2(x_i - x_{i-1})^2}_{\text{from prior}}$$

Where can DP be applied?

Dynamic programming can be applied when there is a linear ordering on the cost function (so that partial minimizations can be computed).

Example Applications:

1. Text processing: String edit distance
2. Speech recognition: Dynamic time warping
3. Computer vision: Stereo correspondence
4. Image manipulation: Image re-targeting
5. Bioinformatics: Gene alignment

Application I: string edit distance

The **edit distance** of two strings, s_1 and s_2 , is the minimum number of single character mutations required to change s_1 into s_2 , where a mutation is one of:

1. substitute a letter (kat \rightarrow cat) cost = 1
2. insert a letter (ct \rightarrow cat) cost = 1
3. delete a letter (caat \rightarrow cat) cost = 1

Example: d(opimizeon, optimization)

```
op imizeon
|| |||||
optimization
|||||
ccicccccsc
```

'c' = copy, cost = 0

d(s_1, s_2) = 2

Complexity

- for two strings of length m and n , exhaustive search has complexity $O(3^{m+n})$
- dynamic programming reduces this to $O(mn)$

Using string edit distance for spelling correction


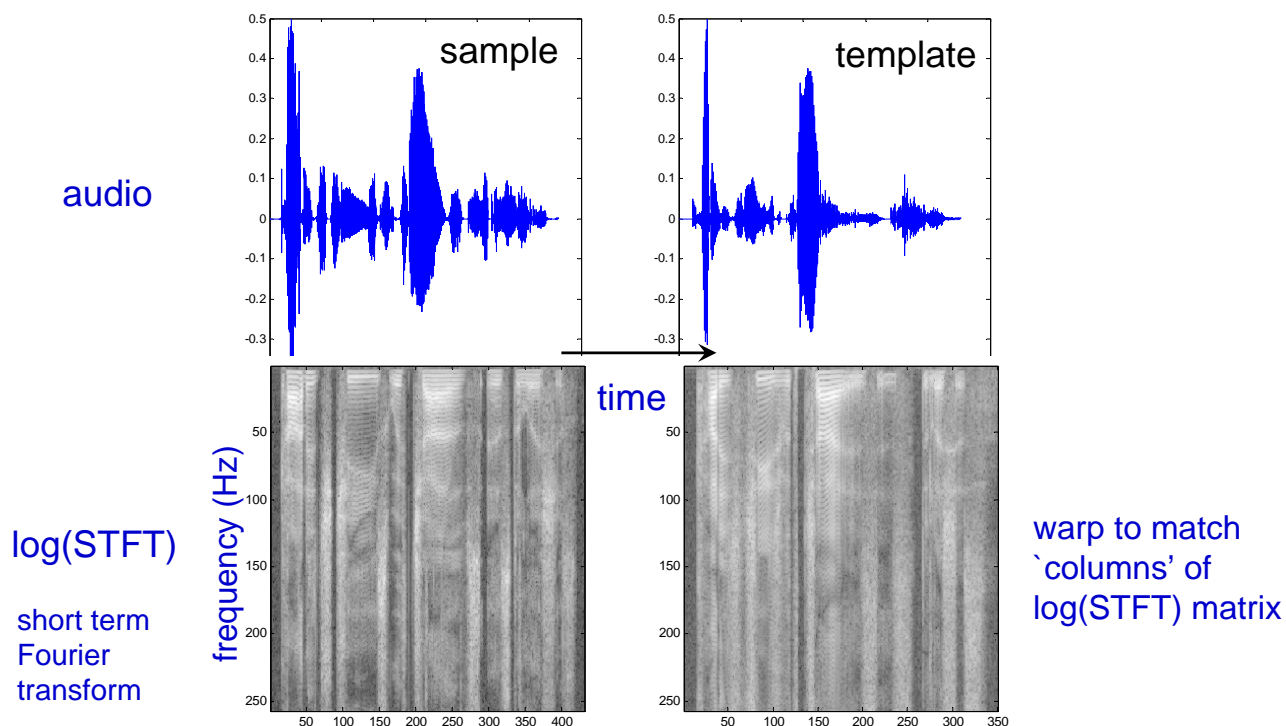
- 1. Check if word w is in the dictionary D
- 2. If it is not, then find the word x in D that minimizes $d(w, x)$
- 3. Suggest x as the corrected spelling for w

Note: step 2 appears to require computing the edit distance to all words in D , but this is not required at run time because edit distance is a **metric**, and this allows efficient search.

Mispelling	Word	ED	Detail
committment	commitment	1	committment \rightarrow commitment
tommorrow	tomorrow	1	tommorrow \rightarrow tomorrow
saftey	safety	2	saftey \rightarrow safty \rightarrow safety

Application II: Dynamic Time Warp (DTW)

Objective: temporal alignment of a sample and template speech pattern



x_i is time shift of i th column

$$f(\mathbf{x}) = \sum_{i=1}^n m_i(x_i) + \sum_{i=2}^n \phi(x_{i-1}, x_i)$$

quality of match \rightarrow cost of allowed moves

$\rightarrow (1, 0)$
 $\downarrow (0, 1)$
 $\searrow (1, 1)$

Application III: stereo correspondence

Objective: compute horizontal displacement for matches between left and right images

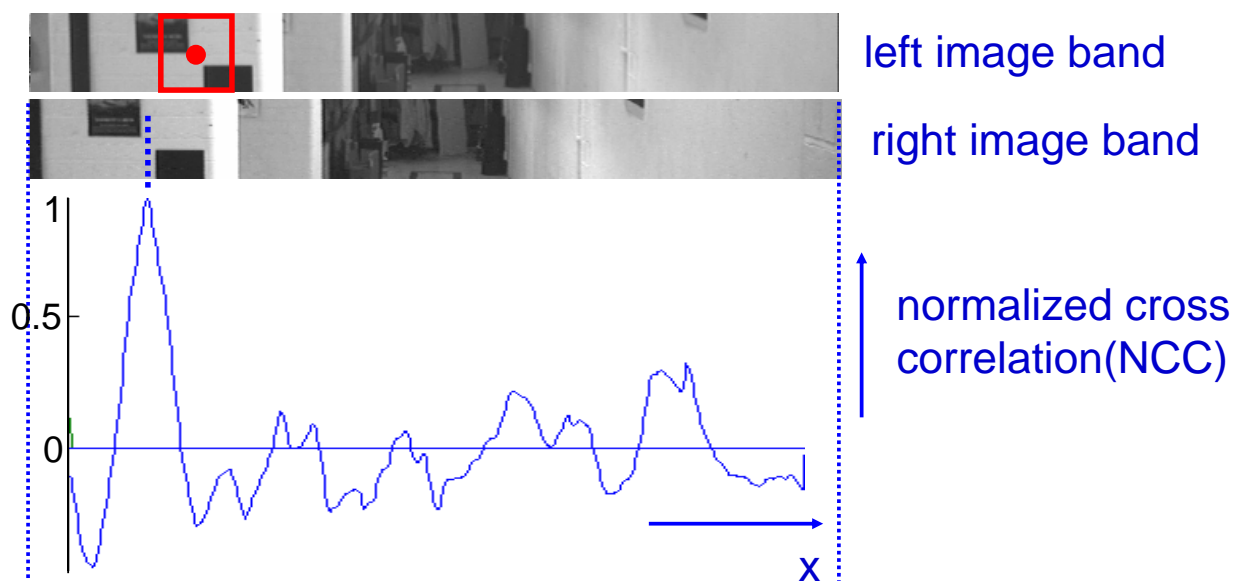


x_i is spatial shift of i th pixel

$$f(\mathbf{x}) = \sum_{i=1}^n m_i(x_i) + \sum_{i=2}^n \phi(x_{i-1}, x_i)$$


quality of match \nearrow uniqueness, smoothness \uparrow

$$m(x) = \alpha(1 - \text{NCC})^2$$



NCC of square image regions at offset (disparity) x

- Arrange the raster intensities on two sides of a grid
- Crossed dashed lines represent potential correspondences
- Curve shows DP solution for shortest path (with cost computed from $f(x)$)



Pentagon example


left image



right image




range map



Real-time application – Background substitution



Left view



Right view

Input



input left view

Results




Background substitution 1



Background substitution 2

Application IV: image re-targeting


- Remove image “seams” for imperceptible aspect ratio change



Seam Carving for Content-Aware Image Retargeting. Avidan and Shamir, SIGGRAPH, San-Diego, 2007



scale




seam
removal



Finding the optimal seam – s

$s \longrightarrow$



$E(I)$



$$E(I) = |\partial I / \partial x| + |\partial I / \partial y| \rightarrow s^* = \arg \min_s E(s)$$